

# Enhancing Reasoning Path Diversity in Commonsense Question Answering

Sagnik Nandi  
23B0905

Harshil Solanki  
23B1016

## Contents

<b>1</b>	<b>Paper Overview</b>	<b>2</b>
1.1	Key Techniques Highlighted:	2
1.2	Model Architectures used:	2
1.3	Sampling Strategies and Self-Consistency:	2
1.4	Limitations and Future Directions:	2
<b>2</b>	<b>Model Overview</b>	<b>2</b>
2.1	Dataset	2
2.2	Pre-trained LLM	3
2.3	Sampling	3
2.4	Output	3
<b>3</b>	<b>Current Progress and Code</b>	<b>3</b>
3.1	Analysis of Strategies	3
3.2	Dataset Analysis	3
3.3	Self-consistency Implementation	4
<b>4</b>	<b>Future Plan of Action</b>	<b>4</b>
4.1	Metrics	4
4.2	Strategies	4
4.3	Timeline	4
<b>5</b>	<b>Code References</b>	<b>4</b>

# 1 Paper Overview

<https://arxiv.org/pdf/2203.11171v4>

**Paper Title:** Self-Consistency Improves Chain of Thought Reasoning in Language Models

## 1.1 Key Techniques Highlighted:

- **Chain-of-Thought (CoT) Prompting:** This involves providing illustrative examples to guide the model's reasoning process.
- **Self-Consistency Decoding:** This method generates diverse reasoning paths and selects the most consistent final answer by marginalizing over them.
- CoT and self-consistency are effective for tasks with fixed answer sets and can be adapted for open-text generation using suitable consistency metrics.

## 1.2 Model Architectures used:

- **Decoder-Only Models (e.g., GPT-3, PaLM, LaMDA):**
  - Well-suited for autoregressive text generation tasks like conversational AI and creative writing.
  - Offer simpler architecture and faster inference times.
- **Encoder-Decoder Models (e.g., T5, BART):**
  - Better suited for sequence-to-sequence tasks such as translation and summarization.
  - Provide stronger input-output alignment but with higher computational demands.

## 1.3 Sampling Strategies and Self-Consistency:

- Sampling techniques (e.g., temperature sampling, top-k truncation) are crucial for generating diverse reasoning paths.
- Self-consistency decoding outperforms traditional methods like sample-and-rank and beam search in terms of accuracy.
- Self-consistency incurs higher computational costs.
- The authors suggest using a limited number of reasoning paths to balance performance and efficiency.

## 1.4 Limitations and Future Directions:

- Self-consistency requires significant computational resources.
- Self-consistency may generate nonsensical reasoning paths.
- Future work could focus on fine-tuning models with supervised data generated through self-consistency to improve single-inference accuracy.
- Grounding models' rationale generation remains an open challenge requiring further research for logical and accurate outputs.

# 2 Model Overview

## 2.1 Dataset

<https://paperswithcode.com/dataset/commonsenseqa>

**CommonsenseQA** is a new multiple-choice question answering dataset that requires different types of commonsense knowledge to predict the correct answers. It contains 12,102 questions with one correct answer and four distractor answers.

## 2.2 Pre-trained LLM

We are using the *together.ai* interface to prompt with SOTA Language models like **Llama 4** and **Deepseek R1**.

## 2.3 Sampling

It first samples a diverse set of reasoning paths instead of only taking the greedy one, and then selects the most consistent answer by marginalizing out the sampled reasoning paths. Self-consistency leverages the intuition that a complex reasoning problem typically admits multiple different ways of thinking leading to its unique correct answer.

## 2.4 Output

Comparison of the self-consistency method with greedy decoding and other baseline methods (e.g., beam search) to highlight its advantages.

# 3 Current Progress and Code

## 3.1 Analysis of Strategies

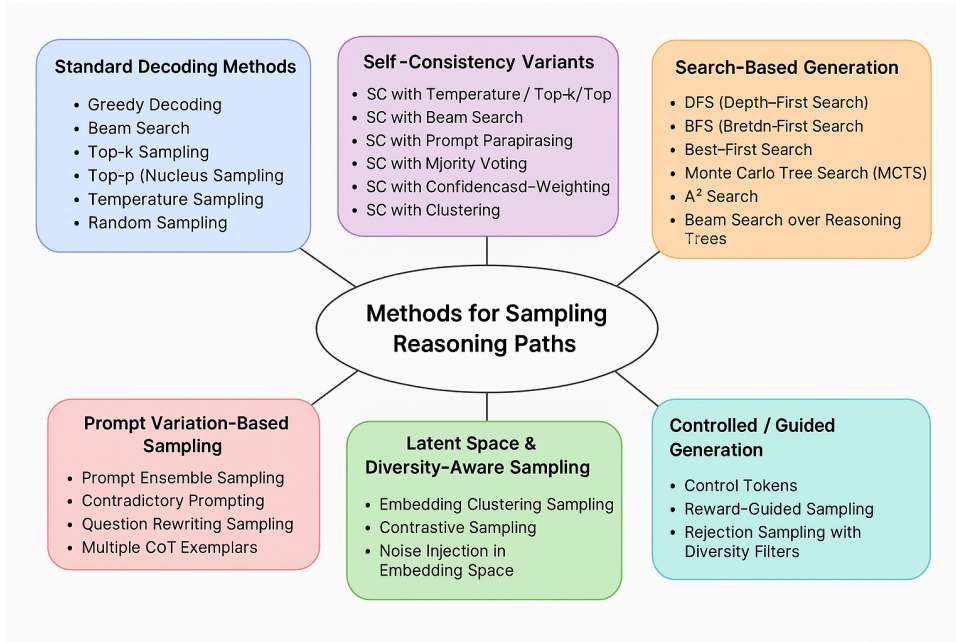


Figure 1: Different Strategies for Sampling

## 3.2 Dataset Analysis

An example of 'train' looks as follows:

```
{
  'id': '075e483d21c29a511267ef62bedc0461',
  'question': 'The sanctions against the school were a punishing blow, and they seemed to _____ the efforts the school had made to change?',
  'question_concept': 'punishing',
  'choices': {
    'label': ['A', 'B', 'C', 'D', 'E'],
    'text': ['ignore', 'enforce', 'authoritarian', 'yell at', 'avoid']
  },
}
```

```
    'answerKey': 'A'
}
```

### 3.3 Self-consistency Implementation

- Ran initial experiments using multiple reasoning paths and beam search
- Choose the most common answer based on 10 responses generated, look at `sc_10_samples.py` in `self_implementation`

Github Link : [https://github.com/HARSHIL2836R/ReasoningLLMs/tree/main/self\\_implementation](https://github.com/HARSHIL2836R/ReasoningLLMs/tree/main/self_implementation)

## 4 Future Plan of Action

### 4.1 Metrics

Provide metrics like:

- Number of unique reasoning paths.
- Path similarity measures (e.g., cosine similarity).
- Diversity scores (e.g., Shannon entropy) to assess the diversity of generated reasoning paths.

### 4.2 Strategies

Basic strategies:

- Sample multiple reasoning paths by **varying temperature** and aggregate answers via majority voting
- **Cluster** the sampled responses based on cosine similarity and select most common answer

Advanced strategies:

- Generate different paths by **varying prompts** like paraphrasing the question
- Model reasoning steps as a graph and use search strategies like **BFS**, **DFS** or **A\*** to discover paths

### 4.3 Timeline

- **April 14–20**: Implement basic strategies and different metric computations and compare their performances
- **April 29-30**: Implement advanced strategies and improving the performance
- **May 1–3**: Final analysis, report writing, and polishing of visualizations and results

## 5 Code References

1. Codebase: [https://github.com/codelion/optillm/blob/main/optillm/self\\_consistency.py](https://github.com/codelion/optillm/blob/main/optillm/self_consistency.py)
2. Drive Link : [https://drive.google.com/drive/folders/1Vzzdk7ZRTuls8XM\\_-3UT0aqaiBEoyV4T?usp=sharing](https://drive.google.com/drive/folders/1Vzzdk7ZRTuls8XM_-3UT0aqaiBEoyV4T?usp=sharing)